GPU MINING BUSINESS MODEL · INVESTOR BRIEF

INFRANEX BT

The decision engine for Bittensor mining: live competition odds, rig-specific dollar projections and tier-flip alerts across all 128 subnets — so capital goes only where a new seat can actually earn.

October 2026

Live network edition · 128 subnets scanned · TAO $254.21 · post-halving emission

Table of Contents

1. Executive Summary 1

2. Bittensor Fundamentals 2

3. TAO Emissions and Distribution 4

4. The GPU Mining Business Model 5

5. Competition Reality: Winner-Take-All 7

6. Infranex BT and the Problem It Solves 10

7. The Feature Suite 12

8. How It Helps Miners Mine 13

9. How the Business Works: Conclusion 15

CHAPTER 1

1. Executive Summary

Bittensor pays real money — roughly 3,600 TAO every day, about $915,000 at the current price — to the people who run the machines behind its AI services. This brief explains that market, the GPU business built on top of it, and Infranex BT: the decision engine that tells a hardware owner where that money can actually be won, and where it is mathematically out of reach.

The opportunity is genuine but brutally selective. Bittensor runs 128 independent AI markets called subnets, each with its own token, its own competition and its own rules for splitting the daily emission. A miner who registers into the wrong subnet does not earn less — the registration fee is burned, the GPU hours are consumed, and the reward share is zero. Our live scan of all 128 subnets on 26 October 2026 graded 104 of them **Don't enter**: the top ten wallets take essentially the entire payout and a new seat earns approximately nothing. Only one subnet is genuinely open, eighteen are fightable, and three seats on the whole network earn a RUN verdict for our reference rig.

128	3,600 TAO	$915k/day	81%
subnets scanned live, every one graded	minted per day, post first halving	emission value at TAO $254.21	of subnets graded Don't enter

What Infranex BT is

Infranex BT is a web application that turns Bittensor's raw chain data into a single honest answer per subnet: can a new seat on this market earn, and what would **your specific rig** make per month if it did. It merges live chain scans with a profitability model on an owned-hardware cost basis, scores every subnet through a conviction engine (55% earning reality, 25% diligence, 20% model confidence), and publishes the result as RUN / WATCH / AVOID verdicts plus a five-tier competition grading from Open field to Don't enter. Dollar projections, month-one win odds and tier-flip alerts are attached to every row.

The product's largest financial line is not a winning pick — it is the avoided loss. A serious GPU entry into the wrong subnet costs roughly $38,000 in hardware plus a non-refundable registration burn plus months of power, against a reward stream that our measurements put at effectively zero for 104 of 128 subnets. The app exists to stop that purchase before it happens, and to route the same capital to the small set of seats where the odds are real. Today that short list pays the reference rig about $881 per month on CPU work (SN05 Numinous) and about $330 per month on each of two GPU seats (SN83 Cathedral, SN92 Beam), with the single best seat on the network — $13,557 per month at 96% month-one odds — reserved by its hardware gate for an H200-class rig.

What this brief covers

- Chapters 2-3: what Bittensor and subnets are, and exactly how TAO emission is minted, halved and distributed.
- Chapter 4: the GPU mining business model — costs, revenue mechanics and worked dollar examples on our rig.
- Chapter 5: the competition reality — why most GPU subnets are winner-take-all, with chain-measured evidence.
- Chapters 6-8: what the app does, its full feature set, and the decision workflow with today's live numbers.
- Chapter 9: how the business works end to end, its revenue logic, growth levers and an honest risk register.

3 RUN · 125 AVOID

today's verdicts for the reference rig across all 128 subnets — the app's value is the 125

CHAPTER 2

2. Bittensor Fundamentals

Bittensor is a proof-of-stake layer-1 blockchain whose entire purpose is to host competitive markets for machine intelligence. It is not a cloud platform and not a token bolted onto a datacenter; the markets themselves are the product, and the chain exists to pay them.

The network in one paragraph

The chain produces one block every twelve seconds and mints TAO against a hard cap of 21 million tokens — a scarcity schedule deliberately modeled on Bitcoin. Every block's emission is distributed to **subnets**: independent, self-contained markets, each run by its own incentive mechanism and each with its own tradeable alpha token. Miners in a subnet produce a digital commodity — inference, model training, data pipelines, agents, storage — and validators score that work every tempo of 360 blocks (about 72 minutes). The scores, weighted by validator stake, decide how the subnet's emission is split among miners. Nothing is promised and nothing is owed: pay follows measured performance.

Subnets: 128 markets under one roof

A subnet is best understood as a company with one product, an automatic payroll, and no moat except performance. Slot 0 is the root network (TAO itself); slots 1 through 127 are live markets covering services from LLM inference (Chutes, SN64) to price prediction, tee-attested compute and agent infrastructure.

Mainnet is capped at 128 active subnet slots today with expansion to 256 planned; the registration cost to claim a slot rose 6.5x by mid-2026, which makes every slot scarce and every incumbent defensive. For a miner, each subnet is a separate job: separate hardware requirements, separate registration burn, separate competition, and a separate answer to "can I win here?" — which is precisely the question the app automates.

Role	What they do	What they earn	What they risk
Miner	Runs hardware and code that performs the subnet's task; registered as a UID	41% slot of the subnet emission, split by consensus score	Registration burn, power, zero pay if scores stay low
Validator	Scores miners' work each tempo; stake-weighted	41% slot, plus a baseline slice to root TAO stakers	Stake slashing risk and the cost of running honest scoring
Subnet owner	Defines the incentive mechanism and rules	Fixed 18% of the subnet's emission	The market itself — a dead subnet earns nothing
Staker / delegator	Stakes TAO or alpha behind validators and subnets	Share of validator emissions, pro-rata	Token price and routing risk across 128 alpha tokens

Table 1 — the four roles of a Bittensor subnet. A GPU owner enters as a miner.

How money actually reaches a miner

Entry requires registering a key into the subnet, which costs a non-refundable recycle-recalibrated fee — demand-priced, and deliberately so: it burns the TAO rather than paying anyone, and it rises when a subnet is attractive. Once in, a miner's pay is decided by consensus: validators submit weight vectors, the Yuma consensus module clips outliers and computes a stake-weighted aggregate, and the miner's share of the 41% miner slot follows that aggregate. Validator trust (vtrust) is itself stake-weighted, which means large stakeholders' opinions dominate the scoreboard. The practical consequence defines this entire business: **you are not selling FLOPS, you are selling credibility to a small committee of staked judges** — and incumbents who have compounded trust for months are hard to displace on hardware alone.

This is also why Bittensor mining is not Bitcoin mining. There is no hash puzzle a faster card wins by default; the work is useful AI labor, the scoreboard is engineered per subnet, and the durable advantage is a code team plus operations plus stake — not the GPU brand on the box. Chapter 5 quantifies exactly how far that advantage goes.

CHAPTER 3

3. TAO Emissions and Distribution

Every twelve seconds the chain mints TAO and routes it through a two-level distribution system: first across 128 subnet markets by market demand, then inside each market by fixed role slots. Understanding both levels is understanding where the money sits.

The spigot: issuance and the halving clock

The block reward is 0.5 TAO — approximately 3,600 TAO per day, worth about $915,000 daily at the live price of $254.21. Bittensor's first halving fired on 14-15 December 2025 when cumulative issuance crossed 10.5 million TAO, cutting the reward from 1 TAO to 0.5 TAO per block; issuance now stands near 11.2 million (53% of the cap) and the next halving triggers at 15.75 million, which at current rates is years away. Each halving also halves every subnet's alpha injection, so scarcity propagates through the whole economy simultaneously. For a miner this is a slowly deflating paycheck in TAO terms — the offset is that demand for subnet services, not the schedule, decides whether any individual seat gets rich.

Milestone	Cumulative issuance	Block reward	Daily emission	Status
Launch (2021)	0	1 TAO	7,200 TAO	done
First halving	10.5M TAO	0.5 TAO	3,600 TAO	fired 14-15 Dec 2025
Second halving	15.75M TAO	0.25 TAO	1,800 TAO	pending (issuance 11.2M)
Third halving	18.375M TAO	0.125 TAO	900 TAO	pending
Hard cap	21M TAO	-	0	asymptotic

Table 2 — the halving schedule. Halvings key off issuance thresholds, not dates.

Level one: routing emission across subnets (dTAO)

Since the Dynamic TAO upgrade of February 2025, the chain does not decide subnet budgets by committee — it runs a market. Every subnet has its own alpha token with a 21 million cap and a TAO/alpha liquidity pool; TAO holders move capital into the pools of the subnets they believe in, and each block's emission is split across subnets in proportion to demand — computed from each pool's EMA price, scaled by burn, and passed through an emission gate that suppresses weak or attacked shares. The root market (plain TAO staking) competes with every subnet for the same capital. The result is Darwinian: a subnet that loses market confidence bleeds emission within days, and a subnet that gains it compounds. Concentration is therefore not a bug of dTAO — it is the mechanism working, and it is why one subnet can legitimately absorb 13% of the entire network's emission while a hundred others share scraps.

Level two: the split inside a subnet

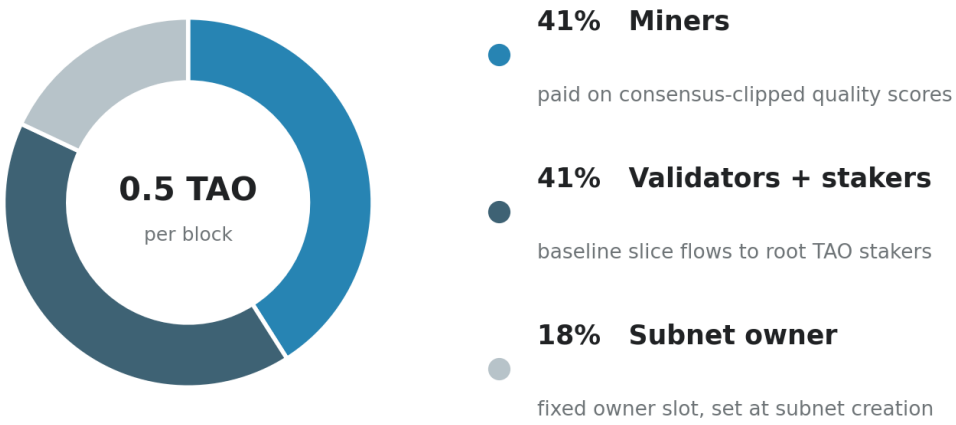


Figure 1 — how one subnet's block emission divides: 41% miners, 41% validators and stakers, 18% owner.

Inside a subnet the split is fixed by hyperparameters: 41% to miners, 41% to validators and their stakers, 18% to the subnet owner. From the validator slot, a baseline slice is redirected to root TAO stakers through the dTAO coupling formula, so plain TAO stakers earn from every subnet proportionally to stake. The miner's 41% is the pot this business competes for. On SN64 — the largest subnet at roughly 13% of all network emission — the miner slot alone receives about 199 TAO per day, near $50,000 daily across the whole miner set; the entire question of GPU mining economics is how many UIDs actually touch that pot, and Chapter 5 answers it with chain data.

≈199 TAO/day

the miner slot of SN64 alone — about $50k/day split across 256 seats, of which 7% earned last epoch

CHAPTER 4

4. The GPU Mining Business Model

Strip the crypto vocabulary away and this is a hardware business: put capital into machines, sell their output into 128 automated markets, and keep the spread between revenue and power. The spread is decided by seat competition — which is why a decision engine, not a bigger power supply, is the scarce asset.

The cost side

The reference rig throughout this brief is an owned 96 GB VRAM workstation (RTX PRO 6000 Blackwell class, TEE-capable) at an owned electricity basis of $0.12/kWh. Three cost lines matter. Capex: a current-generation 96 GB card runs into five figures, and a single wrong-subnet entry effectively strands it — this is the $38,000-class mistake the app exists to prevent. Registration: entering a subnet burns a non-refundable, demand-priced fee, and full subnets recycle constantly so the burn repeats on every re-entry. Power and ops: a loaded GPU draws 400-600 W around the clock — $29-43 per month per card at our basis — plus real engineering time, which is the cost most newcomer models silently omit.

The revenue side

Revenue arrives in TAO and converts at a floating price ($254.21 on the scan date), so every projection carries token risk that no engine can remove. Within the token, two numbers decide the month: the **month-one odds** — our model's estimate that a brand-new seat earns anything at all in its first epoch period, discounted for whale dominance, full subnets and the weeks-long bond ramp — and the **net per month if you earn**, which is expected reward minus power on the active cost basis. A seat is only worth entering when both numbers clear the bar at the same time; the app displays them joined on every row and refuses to quote a payout without the odds that make it real.

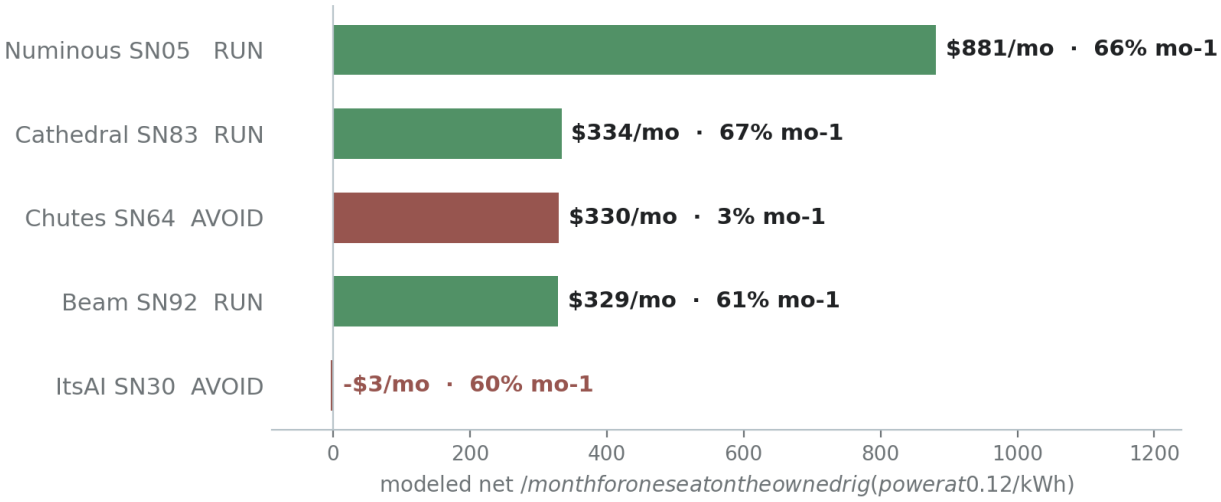


Figure 2 — today's real menu for the 96 GB rig: net $/month per seat with month-one odds. The two AVOID rows show payouts that exist on paper but not in practice.

Worked examples, live numbers

SN05 Numinous is the reference rig's best honest seat: CPU-only work, 66% month-one odds, roughly $881 net per month, conviction 81 — a RUN. SN83 Cathedral and SN92 Beam are the GPU path: about $334 and $329 net per month respectively at 61-67% odds, both RUN, both realistic for a 96 GB card. Contrast SN64 (Chutes): the row still prints $330 per month **if you earn**, but the subnet is full at 256/256 seats, only 7% of seats earned last epoch, and the top-10 wallets take 100% of the pot — month-one odds 3%, verdict AVOID. SN30 (ItsAI) is the subtle trap: 84% of slots earn something, which looks open, but

the top-10% take 99% of the incentive, leaving the median seat $3 per month underwater after power. Payout size without seat odds is marketing; this engine never shows one without the other.

The two-lens rule: seat quality and rig fit

Every opportunity passes two independent gates, and conflating them is the classic newcomer error. The seat lens asks whether a new UID can win at all (Chapter 5's tiers); the rig lens asks whether **this** machine can do the work at an acceptable cost. Today the best seat on the entire network — SN82 Bitcast, an Open field with 96% month-one odds and $13,557 per month net — fails the rig lens for us: it requires an H200-class 141 GB card, so for the reference rig its honest value is exactly zero and its verdict is AVOID. The app is deliberately rude about this. An open seat your rig cannot sit in is not an opportunity, and printing it as one would be selling hope.

$13,557/mo — for an H200 owner.
For this rig: $0.

SN82 Bitcast is the network's best seat; the rig gate is what makes the answer honest

CHAPTER 5

5. Competition Reality: Winner-Take-All

This is the chapter the marketing pages of every mining guide omit. On most GPU subnets the reward distribution is not merely skewed — it is architecturally closed, and the chain data says so plainly.

The honest census

Applying one grading standard to all 128 live subnets — rewarded share of slots, top-10% incentive capture, month-one odds, displacement pressure — produces the census below. One subnet is an Open field where new seats genuinely earn. Eighteen are Fightable: real odds, real warnings. Five are Knife fights where professional farms take most of the pot. And 104 — 81% of the network — are graded Don't enter: a new seat earns approximately zero while still paying the registration burn and the power bill.

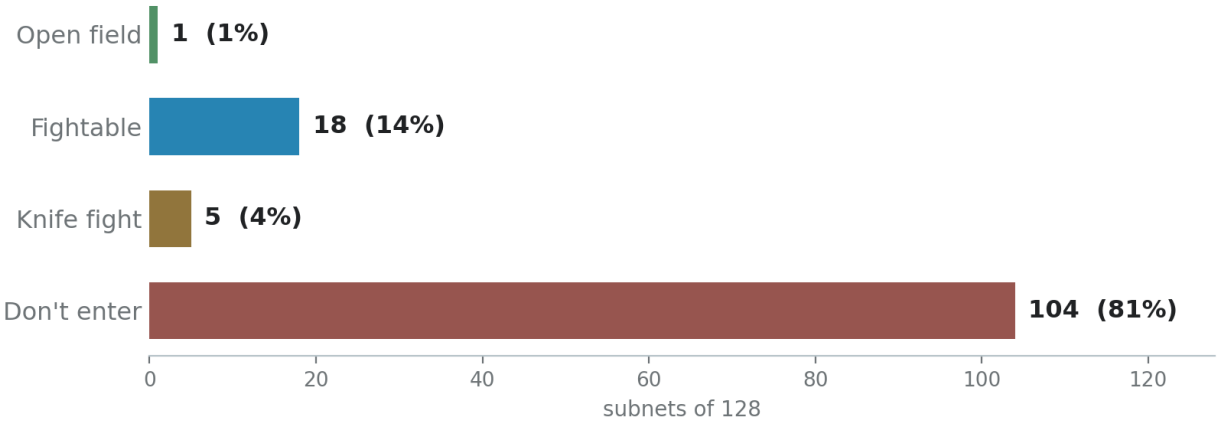


Figure 3 — the competition census across all 128 subnets, live scan of 26 October 2026.

Winner-take-all, measured

The single most predictive number is the share of a subnet's incentive captured by its ten best-paid UIDs. On the flagship GPU subnets that number is not 60 or 70% — it is effectively 100%. Chutes (SN64), the network's largest market and about 13% of all TAO emission, paid 7% of its slots last epoch with its top ten taking the whole pot; its leading operators run Kubernetes farms with dedicated engineering teams and 850 GB to 5 TB of model cache per site. Pareton (SN9) rewarded 0.8% of seats, SayGM (SN27) 0.4%, Targon (SN3) 2.7% — and Targon's minimum hardware line is a rack of 8x H100s with TEE profiles. These are not failing subnets; they are functioning markets with a scale moat, which is exactly why a newcomer with a single GPU should not pay the burn to enter them.

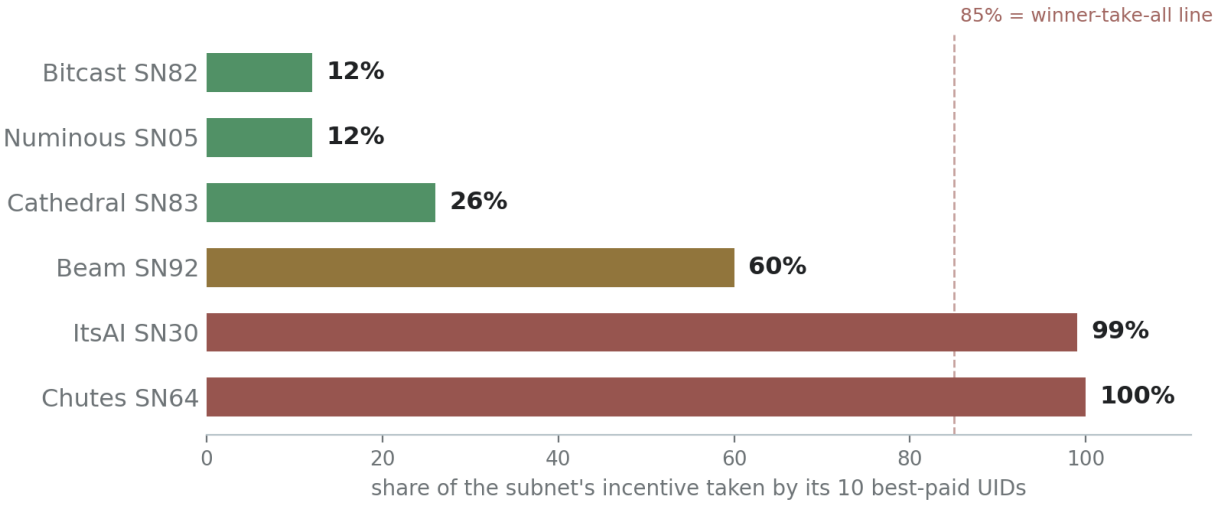


Figure 4 — top-10% incentive capture on case-study subnets. At or above the 85% line, a new seat is fighting the scoreboard itself.

Subnet	Slots earning	Top-10% capture	Min hardware	Verdict
Bitcast SN82	95.5%	12%	H200 141 GB	Open field — rig-gated
Numinous SN05	93.8%	12%	CPU only	RUN
Cathedral SN83	95.7%	26%	GPU 24 GB	RUN
Beam SN92	93.0%	60%	H100 80 GB	RUN
ItsAI SN30	84.4%	99%	GPU 24 GB	AVOID — whale cover
Chutes SN64	7.0%	100%	GPU 24 GB (farms: K8s)	AVOID — knife fight
Pareton SN9	0.8%	100%	GPU 24 GB	AVOID
Targon SN3	2.7%	100%	8x H100, TEE	AVOID

Table 3 — case-study rows from the live scan: earning breadth vs capture concentration.

Why more VRAM does not break the moat

The instinctive response — buy a bigger card — is the most expensive way to be wrong. The moat on the big subnets is three-layered: scale (farms that out-buy any single card), code (teams that optimize against the validator's scoring function weekly), and stake (validators with large delegations whose trust, once earned, is self-reinforcing through the bond ramp). A $38,000 GPU without the engineering stack is a donation to whoever has it. This is why the app's grading keys on distribution spread rather than on hardware class: hardware is table stakes, while the spread tells you whether the table pays anyone but the house.

The SN30 lesson: surface metrics lie, spreads do not

ItsAI (SN30) deserves its own paragraph because it is the trap that catches every newcomer's dashboard. Its 84% slot-earning rate looks like an open field — most seats get something. But the top-10% of UIDs capture 99% of the incentive, so "something" for the median seat rounds to dust: our model puts the median new seat $3 per month underwater after power. Any ranking built on headline numbers — APY, pool size, emission rank — would flag SN30 as attractive. Only a distribution read catches it, and only a distribution read prices it honestly. That is the analytical core Infranex BT sells: not more data, but the two or three numbers that actually decide the outcome.

> *In Bittensor mining, the average is a lie told by the top ten wallets. Read the spread before you pay the burn.*

CHAPTER 6

6. Infranex BT and the Problem It Solves

Everything in chapters 2-5 describes a market that already works. This chapter describes the gap in it: the chain publishes every number a miner needs, and almost nobody can assemble those numbers into an honest answer before their money is already burned.

The problem: information asymmetry with a non-refundable entry fee

Bittensor's data is public — metagraphs, reward spreads, slot occupancy, emission flows — but it is scattered across explorers and APIs, expressed in unfamiliar units, and wrapped in promotional noise from subnet owners whose business model is recruiting miners. The entry fee is a burn: demand-priced, non-refundable, and repeating on every re-entry into a full subnet. Put those together and the newcomer's decision is framed perversely — the moment of maximum enthusiasm is the moment of least information, and the market charges for the experiment. The measurable result is the census in chapter 5: 104 of 128 subnets pay a new seat approximately nothing, yet every one of them collects registrations.

The asymmetry is not symmetrical ignorance, either. The top operators read the spread data daily and position accordingly; validators know their own scoring dynamics; subnet owners know their emission trends. The only participant flying blind is the one writing the largest check for the least information — the hardware owner deciding where to point $38,000. That is the specific asymmetry the app closes, and closing it is the product.

What Infranex BT is

Infranex BT is a web decision engine for Bittensor mining on an owned-hardware basis. It scans all 128 subnets live, merges chain data with a profitability configuration and operator overrides, and produces one graded answer per subnet: a competition tier (Open field, Fightable, Knife fight, Don't enter, No data), a conviction score from 0 to 100, a RUN / WATCH / AVOID verdict, month-one win odds, and a dollar projection specific to the configured rig. Nothing is hand-entered and nothing is cached indefinitely — every verdict traces to a chain read, and when data is stale the app says so on the page instead of pretending otherwise.

The conviction engine: 55 / 25 / 20

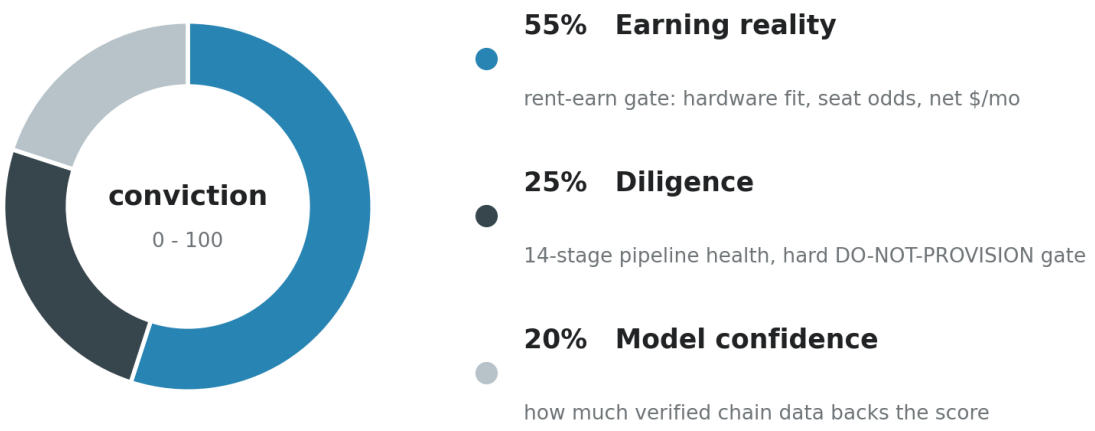


Figure 5 — the conviction engine weights on the owned-rig basis. Earning reality dominates because a seat that never pays invalidates every other signal.

Every subnet is scored by three components. **Earning reality (55%)** is the rent-earn gate: does the rig meet the hardware profile, is a seat even obtainable, what are the odds and the net dollars. **Diligence (25%)** runs a fourteen-stage pipeline over the subnet's operational signals — registration economics, reward history, displacement pressure — and can hard-fail the row outright with a DO-NOT-PROVISION verdict. **Model confidence (20%)** scales the whole score by how much verified chain data backs it, so thin-data subnets cannot ride a lucky average into a RUN. Three hard gates sit above the blend: a failed rig gate, a failed diligence stage, or a below-minimum profit projection each cap the verdict at AVOID regardless of how bright the other numbers look. On today's scan the engine returns 3 RUN, 0 WATCH and 125 AVOID — a distribution that looks harsh precisely because the network is harsh.

The AVOID list is the product

Investors are used to products that sell winners; this one's first deliverable is the 125 rows it refuses to bless. Each AVOID is a prevented experiment: a registration burn not paid, a GPU month not stranded, an engineering weekend not spent porting code into a market whose scoreboard was never going to pay. Priced conservatively at $38,000 of stranded capex plus burn and power per wrong entry, the 104 Don't-enter grades alone represent the difference between a beginner losing their stake and a beginner still holding their capital. The RUN list is small — that is honest — but it is real: today it routes one CPU seat and two GPU seats that together clear roughly $1,540 per month for the reference rig at 61-67% month-one odds each.

125 AVOIDs = 125 prevented $38k mistakes

the app's largest line item is the loss that never happens

CHAPTER 7

7. The Feature Suite

Seven features, one decision. Each exists because a specific question — asked at a specific moment in the miner's workflow — had no honest answer before.

Feature	What it does	The question it answers
Conviction ranking	Scores all 128 subnets 0-100 with hard gates; publishes RUN / WATCH / AVOID	Where do I point capital?
Win-odds panel	Grades every subnet into five competition tiers with month-1 odds, rewarded share and top-10% capture bars	Can I win here at all?
$ projections	Net $/month per subnet on the active rig and power basis, shown only alongside the odds that qualify them	What does MY rig make?
Best-seat banner	Picks the single highest-value enterable seat (Open field + Fightable) for the rig, live	If I enter today, what do I make?
Flip alerts	Diffs every scan against the persisted tier snapshot; surfaces 7-day tier flips as alert chips	When does an impossible seat become possible?
Data honesty notices	Staleness and empty-data states render on-page; UNKNOWN tier never fabricates an answer	Can I trust this row?
Overrides and notes	Operator diligence notes and per-subnet overrides merge into the scoring pipeline	Where do I store my own research?

Table 4 — the feature suite mapped to the decisions it serves.

The win-odds panel and its blunt vocabulary

The panel's five tiers use deliberately unglamorous words. **Open field** means rewarded share is high, whales hold under half the pot, and month-one odds are 50% or better — enter without fear. **Fightable** means real odds but real warnings: displacement required, or a long bond ramp, or whales at half the table. **Knife fight** means professional farms take most of the pot and expect scraps. **Don't enter** means winner-take-all: under 5% of slots earned, or the top-10% capturing 95% or more, or a zero month-one projection — a seat that pays the burn for the privilege of earning nothing. **No data** is the honest shrubbery: where the chain read is incomplete the tier says so instead of guessing. Every row carries its reason line — the actual numbers that produced the grade — so the verdict is auditable at a glance.

Flip alerts: the cheapest signal in the network

Subnet quality is not static. Emission demand moves under dTAO, whale operators migrate, code teams ship, and a subnet graded Don't enter in March can be Fightable by June — that transition is precisely when a new seat is cheapest, before the crowd notices. The flip watcher persists the raw reward-spread tier per subnet on every scan and diffs it against the previous snapshot; qualifying flips surface as a seven-day alert strip in the panel. The same signal in reverse is an exit bell: when a seat's tier degrades from Fightable toward Knife fight, the miner already sitting in it wants to know before the next epoch, not after. One number pair — the tier, and the moment it changes — is the closest thing Bittensor has to a free option.

Dollar projections, honestly scoped

The projection line on each row quotes what the seat pays **if you earn** — expected reward on the active cost basis minus power — and it is always displayed next to the month-one odds that say how likely "if you earn" is. The pair is inseparable by design; a payout without odds is marketing. The best-seat banner collapses the whole panel to the one seat that maximizes rig-fit value among enterable tiers, and the projection math is identical across the app — dashboard, opportunities table and panel share one engine — so the number never disagrees with itself between screens.

CHAPTER 8

8. How It Helps Miners Mine

Features mean nothing until they change a decision. This chapter walks the workflow end to end with today's live scan as the worked example.

The decision workflow

Every cycle the app runs the same pipeline: a chain scan pulls metagraph state for all 128 subnets; the merger joins it with the profitability configuration and any operator overrides; the conviction engine scores and gates each row; the win-odds module grades competition tiers and computes month-one odds; the projection module attaches rig-specific net dollars; and the flip watcher diffs tiers against the persisted snapshot to raise alerts. The miner reads one screen: skip the 104, watch the 18, and choose among the seats that cleared both the seat lens and the rig lens.

Seat (live scan)	Hardware	Conviction	Net $/mo	Mo-1 odds	Verdict
Numinous SN05	CPU only	81	$881	66%	RUN
Cathedral SN83	GPU 24 GB+	75	$334	67%	RUN
Beam SN92	H100 80 GB	71	$329	61%	RUN
Bitcast SN82	H200 141 GB	24	$13,557	96%	AVOID — rig gate
Chutes SN64	GPU 24 GB+	30	$330 at 3% odds	3%	AVOID
ItsAI SN30	GPU 24 GB+	20	-$3	60%	AVOID — whale cover

Table 5 — the short list and the instructive rejects, from the 26 October 2026 scan.

Reading the short list the way an operator would

The CPU-first finding is the headline: the largest honest expected value for this rig is not GPU work at all — Numinous pays $881 per month for CPU-only labor at 66% odds, roughly 2.7x either GPU seat, because CPU markets are less farm-dominated. The GPU path is real but modest: Cathedral and Beam each clear about $330 per month, and a rig with the VRAM headroom can hold both theses in parallel — combined with the CPU seat, the three-RUN portfolio projects about $1,540 per month. Bitcast demonstrates the rig gate doing its job: the best seat on the network is worthless to a 96 GB card, and the engine says so instead of flattering the user. Chutes and ItsAI demonstrate the seat lens: one full, whale-closed market and one whale-covered trap, both priced correctly at approximately zero.

The workflow then becomes maintenance instead of research. The miner holds the RUN seats, checks the flip strip for tier changes, and re-reads the panel when an alert fires — a subnet improving is a candidate entry, a held seat degrading is an exit. Total decision time per week drops from hours of explorer-diving to minutes of reading a screen that has already refused to lie about the other 124 rows.

≈$1,540 / month

combined projection of the three RUN seats for the reference rig, at 61-67% month-one odds each

What the miner actually does, step by step

- Open the panel; read the census and the best-seat banner — thirty seconds for the state of the whole network.

- Filter to Fightable + Open field; for each row read the reason line, the odds and the net dollars together.
- Cross-check the verdict chip: anything AVOID is closed for this rig today, whatever the payout line says.
- Before registering, check the flip strip — a fresh upgrade flip is the cheapest entry; a downgrade on a held seat is the exit bell.
- After entry, track the row weekly: bond ramp progress, rewarded share drift, and the epoch-level odds the model republishes each scan.

CHAPTER 9

9. How the Business Works: Conclusion

The business works because every layer of Bittensor's design routes money through a competition, and every competition creates an information premium for whoever reads it correctly. Infranex BT is that reader, productized.

The complete loop

Follow the dollar once around the system. The chain mints 3,600 TAO a day and routes it by market demand into 128 subnet pools — concentration by design. Inside each pool the emission splits into fixed role slots, and the miner slot is contested by UIDs whose pay follows stake-weighted validator scores — concentration by consensus. The winner-take-all shape that results (chapter 5) makes seat selection, not hardware, the binding constraint on returns — and seat selection is an information problem. Infranex BT sells the missing information layer: live spread analytics, rig-fit gating, dollar projections and flip alerts. Capital that would have burned in 104 closed markets routes instead to the few winnable seats; earnings improve; the operator's data advantage compounds with every scan. The loop closes: better decisions fund better data, which funds better decisions.

Revenue mechanics

The platform's value equation has three lines. First, **avoided loss**: at roughly $38,000 of stranded capex and burn per wrong entry, a single prevented Chutes-class registration repays a subscription many times over — and the app prevents them at a rate of 125 AVOIDs per scan. Second, **earnings uplift**: routing the same rig from a median seat to a RUN seat is the difference between -$3 and $881 per month in today's numbers, a spread the engine quantifies rather than promises. Third, **time**: the weekly decision cycle compresses from hours of manual explorer work to minutes, which is what makes multi-rig and multi-subnet portfolios manageable at all. Monetization follows those lines naturally — a monitoring subscription for the live panel and flip alerts, portfolio tiers for multi-rig operators, and the alert feed as the premium signal for funds that want the earliest read on subnet regime changes.

Growth levers

- Slot expansion 128 to 256: every new subnet re-runs the newcomer's dilemma — and needs the grader.
- dTAO churn: emission re-routing keeps flipping tiers; each flip is a fresh alert event.
- Multi-rig portfolio view: the engine already gates per rig profile; a fleet is the same math times N.
- Staking adjacency: the same spread analytics price alpha-pool risk for pure TAO holders.

The honest risk register

Risk	What it is	How the business absorbs it
TAO price	Revenue is TAO-denominated; $254 today is not a contract	Projections reprice every scan; the decision layer is price-agnostic
dTAO drift	Emission shares move; today's open field can close in weeks	Flip alerts are the productized answer to exactly this risk
Data latency	Scans lag the chain by minutes-to-hours	On-page staleness notices; verdicts cite their scan, never pretend live
Model risk	Odds are estimates from spreads, not guarantees	Odds always shown beside payouts; confidence weighting discounts thin data
Competition	Explorers and subnet marketing pages approximate parts of this	None join spread analytics to rig-fit gating with dollar outputs
Regulatory	Mining-adjacent crypto exposure varies by jurisdiction	The product is information software; custody and token handling stay with the user

Table 6 — what can hurt, and what already handles it.

The one-sentence thesis

> *Bittensor pays $915,000 a day to win competitions whose rules are public but unreadable — Infranex BT is the app that reads them, grades all 128 seats, and alarms when the map changes.*

For the hardware owner, the brief reduces to this: the emission is real, the competition is measured, and the difference between the two is information. The app does not promise that Bittensor mining is easy — today's scan says the opposite, 125 times out of 128. It promises that every seat you consider is one the data can defend, with an alarm attached for the day the answer changes. In a market that charges a burn for every experiment, that is the edge.