

Table of Contents

1. Executive Summary	2
2. Methodology and Data Sources	3
3. The Verified GPU Ranking	4
4. Demand by VRAM Class	6
5. Cross-Checks: Taostat Registry and the INFRANEX Classifier	7
6. The Per-Subnet Evidence Map	8
7. What This Means for GPU Miners	10
8. Coverage Limits and Caveats	11

About this snapshot

This study was produced by the INFRANEX BT DevOps intelligence pipeline as a repo-level audit of GPU requirements across the entire Bittensor subnet registry. Every claim in the chapters that follow is anchored to a verbatim quote from a subnet repository - README files, min_compute.yml hardware declarations, Docker files and dependency manifests fetched at snapshot block 9,036,761 on September 10, 2026, with TAO priced at \$251.03. Chapters 3 and 4 rank GPUs by verified evidence and by VRAM class, Chapter 5 cross-checks the findings against two independent classification systems, Chapter 6 reproduces the full evidence trail, and Chapters 7-8 translate the data into hardware strategy and its limits. The raw JSON datasets (per-subnet evidence and aggregate rankings) accompany this document for independent verification.

1. Executive Summary

This report answers a direct question with repo-level evidence: across all 129 live Bittensor subnets, which GPU is mostly used or required for mining - the H200, the RTX 4090, or something else? We scanned the GitHub repositories linked on-chain for every subnet, pulled each repository's README, its min_compute.yml hardware declaration where one exists, and its Docker and dependency files, then extracted every named GPU model together with the context in which it appears. In total, 108 READMEs (84% of all subnets) and 29 structured min_compute.yml specs were analyzed at snapshot block 9,036,761 with TAO trading at \$251.03.

The headline finding is that **no single GPU model dominates, but the RTX 4090's 24 GB class is the center of gravity for individual GPU miners**, while the A100 is the most frequently named card overall and H100/H200-class datacenter GPUs dominate the highest-paying niches. The A100 appears in 12 subnet repositories (with 9 explicit miner-side requirements, representing 155 TAO/day of miner-side emissions). The RTX 4090 is explicitly named by 6 repositories (4 as a hard miner requirement), and several of the most profitable consumer-GPU subnets publish the 4090 as their reference card. The H200, by contrast, is a niche card in Bittensor mining today: only 4 repositories mention it at all, and its single hard requirement (SN90 KubeTEE) bundles it inside an H100/H200/B200 cluster specification.

A100 Most-named GPU - 12 repos, 9 hard requirements	RTX 4090 Top consumer GPU - 16 subnets in its 24 GB class	H200 Niche - only 4 repos mention it
---	---	--

Measured by reward pool rather than by repository count, the 24 GB class is where the money is: subnets whose explicit minimum fits a 24 GB card (RTX 4090 / 3090 / 5090 class) carry 501 TAO/day of miner-side emissions - about \$3,770,304 per month at current prices - the largest single VRAM bucket in the network. The broad 8-12 GB bucket hosts more subnets (30) but only 340 TAO/day in rewards. The 80 GB+ bucket (A100 80G / H100 / H200) spans 14 subnets and 187 TAO/day, concentrated almost entirely in inference-marketplace and TEE-cluster subnets that rent rather than buy their GPUs. For a solo miner deciding what hardware to acquire, the practical answer is therefore: a 24 GB RTX 4090 remains the most versatile mining card in the Bittensor ecosystem, with an A100/H100 rental strategy as the high-end alternative, and the H200 specifically NOT worth acquiring for Bittensor mining alone at this time.

GPU model	Subnets (required / supported / named)	Miner-side TAO/day behind requirements	USD/month equivalent
A100	9 req / 10 sup / 12 named	154.8	\$1,165,860
RTX 4090	4 req / 5 sup / 6 named	63.7	\$479,409
H100	2 req / 2 sup / 4 named	58.9	\$443,235
A60 / A6000	1 req / 2 sup / 2 named	92.5	\$696,508
H200	1 req / 2 sup / 4 named	54.0	\$406,710

GPU model	Subnets (required / supported / named)	Miner-side TAO/day behind requirements	USD/month equivalent
B200 / GB200	1 req / 1 sup / 3 named	54.0	\$406,710
Apple Silicon	1 req / 1 sup / 1 named	10.2	\$76,574
RTX 5090	1 req / 1 sup / 2 named	3.1	\$23,027

Table 1: Tiered GPU ranking from repository evidence. 'Required' = explicit miner-side requirement; 'supported' = miner-side declared/accepted; 'named' = any mention. Emissions are miner-side TAO/day at snapshot.

2. Methodology and Data Sources

The analysis pipeline ran in four stages. First, we pulled the live subnet registry from the Finney chain through the INFRANEX BT platform's network API at block 9,036,761, capturing 129 subnets with their names, per-subnet emissions and on-chain identity links: 122 of the 129 subnets publish a GitHub URL in their chain identity, and 7 more were recovered through the platform's curated override registry. Second, we fetched each repository's README (trying main and master branches across six common README filenames), plus requirements.txt, pyproject.toml, setup.py, Dockerfile, docker-compose and .env templates, and - critically - min_compute.yml, the Bittensor convention file through which teams declare minimum compute per role. 108 repositories yielded a usable README and 29 shipped a min_compute.yml.

Third, we ran a catalog-based extractor over every text artifact: 30+ canonical GPU patterns (H200, H100, A100, RTX 4090, L40S, A6000, B200 and so on, including variants such as RTX 4090D, 4070 Ti SUPER and MI300X), VRAM minimums ('at least 24 GB', 'min_vram: 80'), and the structured compute_spec.miner.gpu / compute_spec.validator.gpu blocks of min_compute.yml. Every mention was classified into three evidence tiers: **tier 1 - explicit miner-side requirement** (requirement language or a min_compute.yml miner block), **tier 2 - miner-side supported/declared** (mentioned as accepted or declared hardware in a mining context) and **tier 3 - generic mention** (benchmark tables, validator documentation, pricing examples). Fourth, all findings were cross-checked against two independent signals: the taostat community registry's hw_requirements field and the INFRANEX requirements profiler that already powers the platform's DevOps engine.

Two important coverage caveats frame everything that follows. First, 21 of 129 subnets have no readable repository README (they are parked, deprecated, renamed or publish no public code), which removes 23% of miner-side emissions from direct observation; the largest gaps are SN0 (the root network, which is not a minable subnet), SN120 Affine, SN95 Actual and SN97 Albedo, all of which exist but expose no README today. Second, many teams document hardware on external docs sites or in private miner onboarding repos rather than in the public repository, so repo evidence is a floor, not a ceiling: subnets with no GPU mention include some that verifiably require GPUs and are classified by INFRANEX's profiler instead.


Pipeline stage	Coverage	Source / method
Subnets in live registry	129	Finney chain snapshot, block 9,036,761
Subnets with on-chain GitHub link	122	Chain identity registry

Pipeline stage	Coverage	Source / method
READMEs successfully analyzed	108 (84%)	raw.githubusercontent.com, 6 filename variants x 2 branches
min_compute.yml specs found	29	Root + miners/ + docs/ paths
Subnets naming a specific GPU model	27	Catalog of 30+ canonical GPU patterns
Subnets with explicit VRAM minimum	18	README lines or min_compute.yml miner block
Cross-check: taostat registry entries	34/75	Community-maintained hardware registry
Cross-check: INFRANEX classifier	6 GPU classes	Chain metadata + profiler, live via platform API

Table 2: Data coverage at snapshot. Coverage is a floor - external docs and private onboarding repos are invisible to repository scanning.

3. The Verified GPU Ranking

Chart 1 shows the core result: every GPU model named in subnet repositories, with its mentions split by evidence tier. The A100 leads on raw presence - it is the default 'recommended GPU' that subnet teams write into min_compute.yml recommended_gpu fields and README requirement tables, and nearly all of its mentions are genuine miner-side requirements or recommendations rather than passing references. The RTX 4090 follows as the most-named consumer card: five subnets name it as a miner-side requirement or reference configuration, and unlike the A100 these are almost all 'what you actually need to compete' statements rather than aspirational recommendations - for example SN96 Verathos declares required=true with min_vram=24 and recommended_vram=48 anchored on the RTX 4090, and SN53 engy's runbook documents a production 4x RTX 4090 miner.



The datacenter tier tells a sharper story. The H100 and H200 each appear in four repositories, but their hard requirements are concentrated in exactly two high-value subnets: SN90 KubeTEE, whose miners must operate TEE-enabled GPU clusters of H100/H200/B200-class nodes (54 TAO/day of miner-side emissions), and SN40 Ralph, which calibrates its inference marketplace against H100 hardware. The B200/GB200 mentions follow the same pattern - next-generation cluster subnets rather than solo-miner targets. The A6000 deserves special mention: SN12 Compute Horde states it plainly ('requires A6000 - the only GPU supported now'), and SN64 Chutes - the single largest GPU mining opportunity in the network at 92 TAO/day of miner-side emissions - accepts A100, H100 and A6000-class hardware through its node-selector system with a per-GPU VRAM floor of 16 GB. In other words, the highest-value datacenter-GPU mining is marketplace-shaped: the subnet does not hand you a model requirement, it rents whatever VRAM-rich card you bring.




Chart 2 re-weights the same evidence by money: the miner-side emissions of the subnets whose requirements name each card. The picture inverts the count ranking - the A6000 and the marketplace trio (A100/H100 via Chutes and KubeTEE) carry far more TAO per day than their subnet counts suggest, while the 4090's requirement footprint spans several mid-size but steady subnets. For an operator choosing hardware, this is the more honest chart: a 4090 opens five to six durable doors, while an A6000 or rented A100 opens the two largest but most operationally demanding ones.


GPU	Subnets with miner-side requirement	Strongest evidence quote (from repo)
A100	SN12, SN33, SN45, SN49, SN54, SN60...	required=true min_vram=8 rec_vram=24 rec_gpu=NVIDIA A100
RTX 4090	SN9, SN32, SN54, SN96	required=true min_vram=24 rec_vram=24 rec_gpu=NVIDIA RTX 4090
H100	SN40, SN90	- Intel TDX (AMD SEV-SNP, Phase 1) compatible nodes with NVIDIA H100/H200/B200/B300; BIOS + kernel TDX/SGX en...
A60 / A6000	SN64	The most important fields are gpu_count and min_vram_gb_per_gpu. If you wish to include specific GPUs, you ca...

GPU	Subnets with miner-side requirement	Strongest evidence quote (from repo)
H200	SN90	- 8 nodes minimum per cluster (5 control-plane + etcd + worker combined, 3+ dedicated 8-GPU workers per GPU t...
B200 / GB200	SN90	- 8 nodes minimum per cluster (5 control-plane + etcd + worker combined, 3+ dedicated 8-GPU workers per GPU t...
Apple Silicon	SN54	export FLUX_DEVICE="cuda" or mps for Apple Silicon
RTX 5090	SN108	Hardware 8 × RTX 5090 (256 GB) or better CPU only, no GPU

Table 3: Per-GPU requirement evidence with the strongest repository quote. Full quotes for all 27 evidence-bearing subnets appear in Section 6.

4. Demand by VRAM Class

Named models are only half the picture - many subnets specify a VRAM floor rather than a card, and VRAM is what actually gates which hardware you can deploy. We took the strongest explicit miner-side VRAM minimum available for each subnet (min_compute.yml miner block first, then requirement-context README lines), and bucketed the network into five hardware classes. The result is unambiguous: the 24 GB class is the network's center of gravity with 16 subnets and 501 TAO/day (3,770,304 USD/month) of miner-side rewards - both the largest reward pool and a healthy number of distinct opportunities. The 8-12 GB entry class actually hosts more subnets (30, many CPU-optional or light-inference workloads), but pays only 340 TAO/day - broad, cheap competition. The 16 GB middle class (4060 Ti 16G / 4080 / 5080 class) is small but lucrative at 360 TAO/day across 9 subnets, and the 80 GB+ datacenter class rounds out the network with 14 subnets worth 187 TAO/day.



VRAM class	Subnets	Miner-side TAO/day	USD/month	TAO/day per subnet
8-12 GB (3060/4070-class)	30	340.2	\$2,562,092	11.3
16 GB (4060 Ti 16G / 4080 / 5080-class)	9	359.9	\$2,710,397	40.0

VRAM class	Subnets	Miner-side TAO/day	USD/month	TAO/day per subnet
24 GB (RTX 4090 / 3090 / 5090-class)	16	500.6	\$3,770,304	31.3
32-48 GB (A6000 / L40S / A100 40G-class)	3	22.1	\$166,069	7.4
80 GB+ (A100 80G / H100 / H200-class)	14	186.5	\$1,404,669	13.3


Table 4: VRAM-class demand. Subnet counts use the strongest explicit miner-side VRAM minimum; 57 subnets publish no VRAM signal and are excluded from this view.

The per-subnet column of Table 4 is the practical one: the average 24 GB-class opportunity carries roughly 31 TAO/day of miner-side emissions versus 11 TAO/day in the 8-12 GB class - near triple the reward density per opportunity. The 32-48 GB middle (A6000 / L40S class) is conspicuously thin, only three subnets, because in Bittensor's economics you either run consumer-VRAM workloads or jump straight to marketplace datacenter hardware. Note also that VRAM minimums compound with other gates: several 80 GB+ subnets additionally require TEE-capable silicon or cluster-level Kubernetes, which is precisely why their high emissions have not been arbitrated away by an influx of lone miners.

5. Cross-Checks: Taostat Registry and the INFRANEX Classifier

To make sure the repository scan is not an artifact of our own catalog, we compared it against two independent classification systems. The taostat community registry - the oldest public hardware registry for Bittensor - carries hw_requirements entries for 34 of its 75 tracked subnets, and those entries are dominated by links to the same min_compute.yml convention this report extracts directly, which corroborates both the convention's importance and our parsing of it. Where taostat carries free-text requirements they match our findings (for example its entry for the pretraining lineage points at the same min_compute.yml files our scan consumed directly). The registry's coverage ends at roughly the first 24 historically registered subnets, so it serves as a sanity check on the old guard rather than a full-network comparison.

The second cross-check is INFRANEX's own subnet-requirements profiler - the same engine that powers the platform's DevOps Mining Journey. The classifier assigns each subnet a workload category and a representative GPU, interpolating a card when no explicit model is documented. Chart 4 shows its verdict across the network: 65 subnets are classified as runnable on a CPU VPS with no GPU at all, 26 as entry-GPU workloads (12 GB or less), and - most relevant here - **17 subnets are mapped to the RTX 4090 24 GB class carrying 562 TAO/day of miner-side emissions, with 12 more mapped to H100 80 GB**. The classifier's 4090-heavy view is deliberately representative rather than literal: when a subnet says '24 GB VRAM required' without naming a card, the profiler answers the operator's real question - 'what do I rent?' - with the market's default 24 GB card. That the classifier's largest GPU bucket independently lands on the same 24 GB class that our repo evidence ranks first is a strong consistency signal between two very different methodologies.



Two instructive disagreements are worth flagging. First, SN122 CookingTAO - the top net-profit subnet on INFRANEX's opportunity board at roughly \$56,453/month of miner-side emissions - publishes no public README at all: its mining client is distributed through the team's private onboarding, and the profiler classifies it as 'CPU VPS' from chain metadata and category signals alone. That is exactly the kind of subnet where repository evidence goes dark and profiler inference has to carry the load. Second, SN64 Chutes is classified conservatively as an H100-class marketplace, while its repository in fact accepts anything from an A6000 (48 GB) upward through H100/H200/B200 - a reminder that classifier labels are entry heuristics, and the underlying repo text (Section 6) is the authoritative record.

6. The Per-Subnet Evidence Map

This section is the audit trail. Every subnet whose repository names a specific GPU model is listed below with its strongest evidence quote, the file the quote came from, and its tier (1 = explicit requirement, 2 = miner-side supported/declared, 3 = generic mention including validator docs and benchmark tables). Quotes are reproduced verbatim from the repositories, so any figure in this report can be traced back to its source line. Subnets are sorted by miner-side emissions descending, because that is the order in which a miner evaluating where to deploy will encounter them.

Subnet	Name	GPU named	Tier	Verbatim evidence quote (strongest)	Miner TAO/d
SN64	Chutes	A100	1 - requirement	The most important fields are <code>gpu_count</code> and <code>min_vram_gb_per_gpu</code> . If you wish to include specific GPUs, you can do so, w...	92.5
SN90	KubeTEE	B200 / GB200	1 - requirement	- 8 nodes minimum per cluster (5 control-plane + etcd + worker combined, 3+ dedicated 8-GPU workers per GPU type) — all...	54.0
SN9	iota	RTX 4090	1 - requirement	1. Cuda GPU with = 16GB VRAM (RTX 4090, for example).	35.4

Subnet	Name	GPU named	Tier	Verbatim evidence quote (strongest)	Miner TAO/d
SN53	engy	RTX 4090	2 - supported	miner/ folder on its PYTHONPATH (the runbook has a working 4 × RTX 4090	30.7
SN49	Nepher Robotics	A100	1 - requirement	Requires NVIDIA GPU (A100+ recommended), Isaac Sim 5.1, Isaac Lab 2.3.0, Docker + NVIDIA Container Toolkit.	16.3
SN34	BitMind	A100	3 - mention	required=true min_vram=80 rec_vram=80 rec_gpu=NVIDIA A100 80GB PCIE	14.9
SN96	Verathos	RTX 4090	1 - requirement	required=true min_vram=24 rec_vram=48 rec_gpu=NVIDIA RTX 4090	14.3
SN17	404—GEN	H200	2 - supported	2. Pod deployment: Deploy the image on a GPU pod via the batch generation API, using the hardware the miner declared in...	12.8
SN10	Pareton	H200	3 - mention	SGLang campaign for Qwen/Qwen3.8-27B-FP8 on one H200.	12.0
SN83	CliqueAI	A100	1 - requirement	required=false min_vram=8 rec_vram=24 rec_gpu=NVIDIA A100	11.6
SN91	cascade	L40S	3 - mention	chain.toml ships with mainnet values baked in: netuid 91, the L40S GPU pin,	10.6
SN54	Yanez	Apple Silicon	1 - requirement	export FLUX_DEVICE="cuda" or mps for Apple Silicon	10.2
SN104	TAOstatus	A100	2 - supported	Three miners, a 100-unit reward pool for one epoch:	7.7
SN63	Enigma	RTX PRO 6000	2 - supported	Factor large semiprimes to break RSA encryption. Solutions run in Docker containers on validator hardware (NVIDIA RTX P...	7.5
SN12	Compute Horde	A60 / A6000	2 - supported	1. trusted miner (requires A6000 - the only GPU supported now) for cross-validation	6.1
SN33	ReadyAI	A100	1 - requirement	required=true min_vram=8 rec_vram=24 rec_gpu=NVIDIA A100	5.7
SN60	Bitsec.ai	A100	1 - requirement	required=true min_vram=8 rec_vram=24 rec_gpu=NVIDIA A100	5.3
SN40	Ralph	H100	1 - requirement	H100 calibration (matmul) 0.512 ms	4.8
SN74	Gittensor	RTX 5090	3 - mention	Serving (compute): to pay compute miners a validator needs its own RTX 5090 running the reference runtime —	4.4
SN125	Refinery	B200 / GB200	3 - mention	Compute One B200; at most 20 hours	4.1
SN117	glyph	RTX 4090	3 - mention	Default eval path is local Docker on an RTX 4090 — every validator running GPU codecs must	3.9
SN29	horfloat	B200 / GB200	3 - mention	GPU requirements: NVLink/SXM 8-GPU pod (8x H200, 8x B200, or 8x B300 are Tier A; 8x H100 is Tier B fallback), 400 GB st...	3.9
SN113	TensorUSD	A100	1 - requirement	required=false min_vram=8 rec_vram=24 rec_gpu=NVIDIA A100	3.8
SN32	ItsAI	RTX 4090	1 - requirement	required=true min_vram=24 rec_vram=24 rec_gpu=NVIDIA RTX 4090	3.8
SN45	AlphaRidge.ai	A100	1 - requirement	required=true min_vram=8 rec_vram=24 rec_gpu=NVIDIA A100	3.3

Subnet	Name	GPU named	Tier	Verbatim evidence quote (strongest)	Miner TAO/d
SN108	Prometheon	RTX 5090	1 - requirement	Hardware 8 × RTX 5090 (256 GB) or better CPU only, no GPU	3.1
SN72	StreetVision by NATI	GTX 1660 family	1 - requirement	required=true min_vram=4 rec_vram=8 rec_gpu=NVIDIA GTX 1660 Super	3.0

Table 5: All 27 subnets whose repositories name a specific GPU model, sorted by miner-side emissions. Where a repository names several GPUs the strongest-tier one is shown; full per-subnet detail is in the accompanying JSON dataset.

6.1 Subnets where repository evidence is absent

For completeness, the following higher-emission subnets expose no readable repository README and therefore cannot contribute evidence here: SN0 (root network - not minable), SN120 Affine (134.0 TAO/day, repository exists but ships no README), SN95 Actual (105.8), SN97 Albedo (59.4), SN110 Green Compute (31.9) and SN122 CookingTAO (11.5, private mining client). A further group - SN16, SN30, SN31, SN42, SN73, SN87, SN99, SN109, SN112 - have malformed or parked chain identities ('0x' placeholders or dead links) and are effectively dormant. None of these change the ranking conclusions, but SN120 and SN95 in particular are large enough that their true requirements deserve a manual check through their official documentation before committing hardware, and the INFRANEX profiler classifications quoted in Section 5 are the best standing estimates for them.

7. What This Means for GPU Miners

If you own or plan to acquire a single consumer GPU for Bittensor mining, the evidence points to the **RTX 4090 (24 GB)** as the most versatile card in the network: it is the named reference card in five subnets' miner requirements, it satisfies the entire 24 GB VRAM class (16 subnets, 501 TAO/day of miner-side rewards), and it clears the 16 GB class (9 subnets) as well. The RTX 5090 appears as an accepted alternative where noted (SN108), but subnets do not yet require it anywhere, so paying the 5090 premium purely for Bittensor is not supported by today's evidence. Cards below 16 GB are viable in the 8-12 GB class but face the thinnest reward density per subnet, and several of those subnets are CPU-optional, meaning GPU-less competitors can dilute them.

If you rent datacenter capacity instead, target the two marketplace-shaped giants: SN64 Chutes (92 TAO/day of miner-side emissions, A6000/H100/H200 accepted with 16 GB per-GPU minimums) and SN90 KubeTEE (H100/H200/B200 TEE clusters), plus the nine A100-anchored subnets from Table 3. The H200 question deserves a direct answer, since it motivated this study: **the H200 is neither required nor forbidden anywhere on its own merits** - it appears only inside KubeTEE's H100-or-newer cluster menu and as a passing option elsewhere, so an H200 purchase justified by Bittensor alone would rest on one subnet's co-equal option. An H100 or A100 rental unlocks strictly more subnets today at lower cost.

Your hardware	Top target subnets from the evidence	Reward pool unlocked
RTX 4090 / 24 GB class	iota (SN9), ItsAI (SN32), engy (SN53), Yanez (SN54), Verathos (SN96)	24 GB class: 16 subnets, 500.6 TAO/d

Your hardware	Top target subnets from the evidence	Reward pool unlocked
RTX 5090	SN108 plus any 24 GB-class subnet	Accepted in 2 repos; required by none
A6000 48 GB	Compute Horde (SN12), Chutes (SN64)	Chutes alone carries 92.5 TAO/d
A100 / H100 80 GB	Chutes (SN64), KubeTEE (SN90), + 9 A100-anchored subnets	80 GB+ class: 14 subnets, 186.5 TAO/d
H200	KubeTEE cluster menu only (with H100/B200)	No solo-miner requirement anywhere

Table 6: Hardware-to-subnet strategy map derived from the verified evidence.

Operationally, fold this research into the INFRANEX Mining Journey: in step 1 (choose what to mine) the Opportunities board already ranks subnets by your profitability target; use Section 6's evidence table to check the named-GPU requirement before step 2 (rent a GPU host), and let the DevOps Engine's per-subnet requirements profiler pull the authoritative, freshly-cached specification for your chosen subnet before you commit to a rental. The profiler's per-subnet pull is the live counterpart of this report's snapshot - emissions, prices and requirements all drift week to week, and the ten-step validation pipeline will verify the host you actually rented against the subnet's real driver, CUDA and VRAM gates before any miner is installed.

8. Coverage Limits and Caveats

Four limitations bound these conclusions, and none of them is silent - each can be checked or re-measured. First, coverage: 108 of 129 READMEs is strong by count, but the unreadable tail holds 23% of miner-side emissions, concentrated in SN0, SN120, SN95 and SN97; any claim about those specific subnets in this report is explicitly marked as inference. Second, documentation culture: repository READMEs undersell GPU requirements because teams maintain setup guides on docs sites, Discord and private onboarding repos; the tier system exists precisely so that soft signals are never presented as hard requirements, and the classifier cross-check in Section 5 exists so that silence in repositories is not mistaken for absence of GPUs.

Third, extraction judgment: regex catalogs cannot perfectly judge intent, which is why every claim ships with its verbatim quote - the strongest quotes were manually reviewed during analysis, and borderline cases (benchmark tables, marketplace node-selectors, miner-declared hardware) are tiered conservatively; where a tier genuinely matters to a decision, read the quote, not the tier. Fourth, time: this is a snapshot at block 9,036,761 (September 10, 2026, TAO at \$251.03); subnet requirements change with releases, emissions drift daily, and the RTX 5090's near-absence from requirements will presumably age as the installed base grows. Re-running the pipeline is a five-minute exercise on the INFRANEX platform - treat this document as the baseline that future snapshots are diffed against, not as a permanent truth.