

Table of Contents

1. Verdict Up Front	1
2. Scope, Method and Data Sources	2
3. Provider-by-Provider Review	2
Vast.ai — deepest market, weakest guarantee	2
RunPod — the managed middle that keeps pods alive	3
Akash Network — decentralized bids, mid-market pricing	3
Lambda Labs — the enterprise fallback	3
The rest of the field — cheap until they are not	3
4. Runtime and Reliability: The Decisive Metric	4
5. Top GPU Performers and Price Benchmarks	5
6. Provider Fit for Subnet Mining	6
7. Operator Action Checklist	7
8. Limitations and Sources	8

1. Verdict Up Front

RunPod is the best default for running rented GPUs when runtime is the metric that matters. It is the only mainstream provider in this comparison that publishes a documented uptime SLA (99% on Secure Cloud), and it is the only one whose failures are priced into a managed platform rather than pushed onto individual marketplace hosts. For a Bittensor miner this is the whole game: a crashed miner earns zero emissions while the invoice keeps running, so the provider that keeps a pod alive through host maintenance, driver faults and noisy-neighbour churn is worth a premium per hour.

#	Provider / tier	Grade	H100 \$/hr (typ.)	Uptime posture	Platform status
1	RunPod — Secure Cloud	A	2.89	99% documented SLA; managed infra	Offers + real rental adapter
2	Lambda Labs	A-	2.86	Consistent; enterprise support	Offers only (rent via RunPod)
3	RunPod — Community Cloud	B+	1.99	No formal SLA; multi-tenant	Same RunPod key
4	Akash Network	B	~3.37 median	Provider-dependent; no SLA	Public bid offers
5	Vast.ai — verified hosts	B-	1.73–2.16	No SLA; host ratings + benchmarks	Offers + real rental adapter
6	Vast.ai — unverified / spot	C+	0.90–1.20	Interruptible; +20–40% effective cost	Same Vast.ai key
7	TensorDock / Hyperstack / Clore	C+	2.25 claimed	Thin track records; TensorDock issues	Not integrated

Table 1: Provider ranking for continuous rented-GPU mining (observed Sep 2026)

99% vs none

RunPod Secure Cloud publishes a 99% uptime SLA — the only documented runtime guarantee in this field. Everyone else asks you to trust the host.

Vast.ai is the price floor, not the reliability leader. Its marketplace routinely lists H100-class capacity below \$1.20 per hour and its median SXM rate undercuts every managed platform, but there is no SLA and reliability depends entirely on which host you pick. Market-observed comparisons put the effective cost of unverified Vast hosts 20–40% above the sticker once downtime and re-runs are priced in. Lambda Labs is the enterprise fallback — consistent uptime and real support at a modest premium — while Akash Network sits in the middle as a decentralized, bid-priced market whose quality varies by provider. TensorDock and the newer entrants (Hyperstack, Clore.ai, Fluence) are cheap on paper but have the thinnest reliability track records, with TensorDock reporting serious availability problems since late 2025.

2. Scope, Method and Data Sources

This brief answers one operator question: which provider keeps rented GPUs earning the longest, and which specific GPU models deliver the most mining value per hour across those providers. Five providers receive a full review — Vast.ai, RunPod, Akash Network, Lambda Labs, and a combined tier of alternatives (TensorDock, Hyperstack, Clore.ai, Fluence). Each is scored on five dimensions: uptime posture and SLA honesty, headline and effective price, offer depth across GPU classes, interrupt behavior on spot or interruptible inventory, and practical fit for Bittensor subnet workloads.

The runtime-first weighting is deliberate. Proof-of-learning and inference subnets pay for continuous availability, not for peak throughput, so a provider with a slightly higher sticker price but a hard SLA beats a cheaper marketplace host that reboots weekly. Price observations are medians and typical advertised rates drawn from provider pricing pages, live platform catalog adapters (the platform's provider registry integrates RunPod, Vast.ai, Lambda and Akash directly) and third-party market comparisons published between June and September 2026. All figures are observed market rates, not quotes, and Section 8 lists the source set with dates so every number can be re-checked before a large lease.

3. Provider-by-Provider Review

Each provider below is reviewed on the five scoring dimensions, with the numbers that an operator actually negotiates against: H100-class on-demand rates, SLA posture, and what the platform's own adapters can automate.

Vast.ai — deepest market, weakest guarantee

Vast.ai is a two-sided marketplace where third-party hosts set their own terms, which produces the widest price range and the widest quality range in the industry. H100 listings start under \$1.00 per hour, the SXM median sits around \$2.16, and H200 capacity is the deepest of any platform at a median near \$4.65. Reliability tooling is real but self-service: machine reliability ratings, verified benchmarks and a scoring system help you avoid the worst hosts, yet nothing prevents a host from taking a machine down for maintenance. Interruptible instances priced by bidding reward operators who build auto-restart into their miners. In the platform, Vast.ai offers feed both the GPU and CPU catalogs and the rental adapter provisions on-demand instances with the SSH key injected — spot inventory must be managed manually.

RunPod — the managed middle that keeps pods alive

RunPod operates the infrastructure it sells, and that shows up in the failure model: Secure Cloud carries a documented 99% uptime SLA, publicly documented compliance, and a standardized path through Pods, Serverless and Clusters. H100 PCIe runs \$1.99 per hour on Community Cloud and \$2.89 on Secure Cloud; A100 PCIe is \$1.19 and \$1.59 respectively. Community Cloud is the value play for rerun-tolerant work, while Secure Cloud is where a miner that must not go down belongs. The platform's RunPod adapter is the most complete integration available: live pricing plus real rentals, so a bundle picked in the deploy wizard becomes a running pod with monitoring attached.

Akash Network — decentralized bids, mid-market pricing

Akash aggregates GPU bids from a distributed provider set, and its public pricing needs no API key — H100 on-demand leases have ranged from \$2 to \$7 with a median near \$3.37 in recent market data. The upside is censorship resistance, transparent bid aggregation and occasional underpriced capacity from providers building reputation. The downside is structural: leases inherit whatever reliability the individual provider delivers, there is no SLA to appeal to, and the platform's Akash feed is offer-only with rentals handled through other adapters. Akash is best treated as a secondary market to check before accepting a managed-platform rate, not as the home of an uptime-critical miner.

Lambda Labs — the enterprise fallback

Lambda sells on consistency rather than price: H100 PCIe on-demand around \$2.86 per hour, 1-Click Clusters for multi-GPU work, no egress fees, and the most enterprise-grade support posture in the comparison. Users report reliable access to A100 and H100 inventory for sustained training runs, and uptime is consistently described as solid across third-party comparisons. The premium over RunPod Community is small — often a few cents per GPU-hour — and buys predictable capacity instead of marketplace variance. The platform pulls Lambda's live on-demand pricing into the catalog, but the rental adapter is not live, so Lambda capacity is priced here and provisioned via RunPod.

The rest of the field — cheap until they are not

TensorDock has advertised some of the lowest H100 SXM rates in the industry (claimed \$2.25) but has had serious availability problems since late 2025, which converts a good sticker price into a bad effective one exactly when capacity is scarce. Hyperstack targets mid-tier European capacity with competitive on-demand rates and a smaller track record. Clore.ai and Fluence represent the decentralized wave behind Akash: real savings on individual hosts, marketplace-style reliability variance, and rating systems that push the diligence work onto the operator. None are integrated into the platform's adapters today, so using them means manual key management and manual monitoring outside the dashboard.


4. Runtime and Reliability: The Decisive Metric

Mining economics turn sticker price into effective price through uptime. A miner that is offline produces no emissions but still accrues rent, so the true hourly cost of a GPU is the sticker divided by the fraction of time it actually works. At 99.5% uptime the adjustment is cosmetic; at 92% — typical for unverified marketplace hosts that reboot for maintenance or lose machines to higher bidders — it silently erases most of the discount. This is why the SLA question is not administrative paperwork but the core of the price comparison.

Provider / tier	Sticker \$/hr	Assumed uptime	Effective \$/hr	What the gap buys you
Vast.ai — unverified floor	1.20	92%	1.30	Cheapest entry; you own the restart problem
RunPod — Community	1.99	97%	2.05	Managed multi-tenant; rerun-tolerant work
Vast.ai — verified median	2.16	97%	2.23	Host ratings + benchmarks, still no SLA
TensorDock	2.25	92%	2.45	Low sticker; availability track record poor
Lambda Labs	2.86	99.5%	2.87	Enterprise consistency, real support
RunPod — Secure Cloud	2.89	99.5%	2.90	The only documented 99% SLA
Akash — bid median	3.37	97%	3.47	Decentralized access; provider lottery

Table 2: Sticker price vs downtime-adjusted effective price, H100 class (observed Sep 2026)

Two readings fall out of the table. First, once downtime is priced in, the cheapest unverified Vast floor (~\$1.30 effective) and RunPod Secure (~\$2.90 effective) are only about \$1.60 per GPU-hour apart — before counting the operator hours spent babysitting restarts. Second, the middle of the market compresses hard: TensorDock's low sticker lands above Lambda's effective price, and Akash's bid median is the most expensive entry once its provider-dependent uptime is assumed. Interruptible spot inventory deserves a separate note: at 30–60% discounts it is genuinely profitable for rerun-tolerant subnets, but only with automated restart and checkpointing, because interruptions are a when, not an if.



Ra nk	GPU	VRA M	Typical \$/hr	Sweet-spot provider	Mining fit
1	H100 SXM	80 GB	1.73–2.89	RunPod Secure; Vast verified	Emissions-per-dollar workhorse; baseline on high-value GPU subnets
2	H200	141 GB	3.68–4.65	Vast.ai (deepest supply)	Memory-bound training and long-context inference subnets
3	RTX 4090	24 GB	0.40–0.55	Vast.ai (consumer fleet)	Efficiency king for lightweight inference and 4090-native subnets
4	A100 PCIe	80 GB	1.19–1.59	RunPod Community	Price/performance floor for LLM inference loops
5	B300	288 GB	~8.69–8.75	Vast.ai; Lambda	Frontier class; only where emissions clearly justify the rate

Table 3: Top GPU performers for rented mining, ranked by value per GPU-hour (Sep 2026)

Ranking logic is deliberate. The H100 SXM takes the top slot not because it is the fastest card but because it clears the admission bar of the most subnets while staying inside a price band that effective-cost math can still justify. The H200 outranks the 4090 only where memory is the binding constraint; for pure tokens-per-dollar on light workloads, the 4090's sub-\$0.55 rate is unmatched. The A100 survives as the budget inference floor, and B300-class silicon is a frontier specialist — rent it only when a subnet's emissions demonstrably scale with frontier compute, otherwise the \$8.75 median rate pays for capacity the workload never uses.

6. Provider Fit for Subnet Mining

The right routing depends on what a subnet does to you when the GPU disappears. Continuous-scoring subnets punish downtime twice — missed validation windows plus a cold-start penalty on re-registration — so they belong on SLA-backed capacity. Rerun-tolerant batch workloads can harvest interruptible discounts safely as long as restart is automated. The table below maps the platform's current top GPU-scored opportunities onto that logic; treat the subnet assignments as examples to re-check against the live dashboard, since opportunity scores rotate weekly.

Workload pattern	Example platform targets	Route to	GPU class	Why
Continuous uptime scoring	SN5 Hone, SN41 Almanac, SN123 MANTIS	RunPod Secure	H100 SXM	SLA-backed runtime; restarts cost emissions
Interruptible batch	SN61 RedTeam, SN67 Harnyx	Vast.ai spot	RTX 4090 / H100	30–60% cheaper; auto-restart masks interrupts
Budget inference loops	Light GPU / mixed catalogs	Akash bids	A100	Bid-priced capacity; verify provider history first
Enterprise lease, fixed term	Any high-emission GPU subnet	Lambda	H100 PCIe	Consistent capacity for predictable monthly cost

Table 4: Routing guide — workload pattern to provider and GPU class

Two platform-specific notes sharpen the routing. First, the rental adapters that the dashboard can automate today are RunPod and Vast.ai (on-demand), which conveniently covers both ends of the reliability spectrum; Lambda and Akash offers are visible but their capacity must be provisioned manually. Second, the GPU catalog's spot flags and availability grades (available / limited / scarce) are exactly the fields to filter on before leasing: scarce availability plus no SLA is the combination that produces multi-day outages, and it is visible in the catalog before the wallet ever pays.

7. Operator Action Checklist

Seven steps convert this analysis into a running configuration. Each is small; together they keep the rigs earning.

- 1. Point uptime-critical miners at RunPod Secure Cloud.** Configure the RunPod API key in the platform key vault and treat the 2.89 H100 rate as the ceiling for continuous subnets — the SLA is the product being bought.
- 2. Keep Vast.ai for interruptible bids only.** Use verified hosts with high reliability ratings, on-demand for anything that cannot survive a restart, and bid-priced spot only with automated restart and checkpointing in the miner.
- 3. Price every lease in effective dollars, not sticker.** Divide the hourly rate by the host's realistic uptime before comparing providers; the +20–40% unverified-host premium is invisible on the listing page.
- 4. Verify the H100 median before long leases.** The market has moved more than 20% in twelve months; re-check current medians (\$2.16 SXM on Vast, \$1.99/\$2.89 RunPod) weekly before committing to multi-week capacity.
- 5. Filter the catalog on availability before price.** Scarce-availability offers with no SLA are the source of the worst outages; a limited-availability Secure Cloud pod out-earns a scarce marketplace bargain.

- 6. Use Akash and Lambda as leverage, not as home.** Check Akash bid medians and Lambda on-demand rates before accepting a managed rate; use the spread to negotiate or to route rerun-tolerant work.
- 7. Re-run this comparison quarterly.** Provider reliability is a moving target — TensorDock deteriorated within a year of being a price leader, and decentralized entrants improve monthly.

The one-sentence answer to the original question: for rented GPUs with good runtime, RunPod — Secure Cloud for continuous miners, Community for the tolerant rest — is the best provider in the field, with Vast.ai as the disciplined operator's discount arm, Lambda as the enterprise fallback, and Akash as a secondary market worth checking before any large lease.

8. Limitations and Sources

Three limitations bound this analysis. Pricing on marketplaces moves continuously; every dollar figure here is an observed median or advertised rate from June to September 2026 and should be re-verified against live catalogs before large leases. Marketplace reliability is host-level, so provider-level grades describe the distribution, not every machine on it — a well-chosen verified Vast host can and does outperform a bad managed node. Finally, platform adapter status determines what can be automated from the dashboard: RunPod and Vast.ai support live offers plus real rentals, Lambda and Akash expose live pricing only, and the remaining providers are outside the adapter set entirely.

Primary sources consulted: provider pricing and documentation pages for RunPod, Vast.ai, Lambda Labs and Akash Network (accessed September 2026); the platform's own provider registry and GPU catalog offer adapters, which pull live H100/H200/A100/4090 pricing from the integrated providers; third-party comparisons including Tech Insider's RunPod vs Lambda vs Vast.ai pricing comparison (Aug 2026), Spheron's RunPod vs Vast.ai cost analysis (Jul 2026), Clore.ai's cheapest-GPU-cloud comparison (Jan 2026), and Akash Network's H100 rental price survey (Aug 2026). Where sources disagreed, the more conservative uptime assumption was used, and ranges are shown instead of false precision.